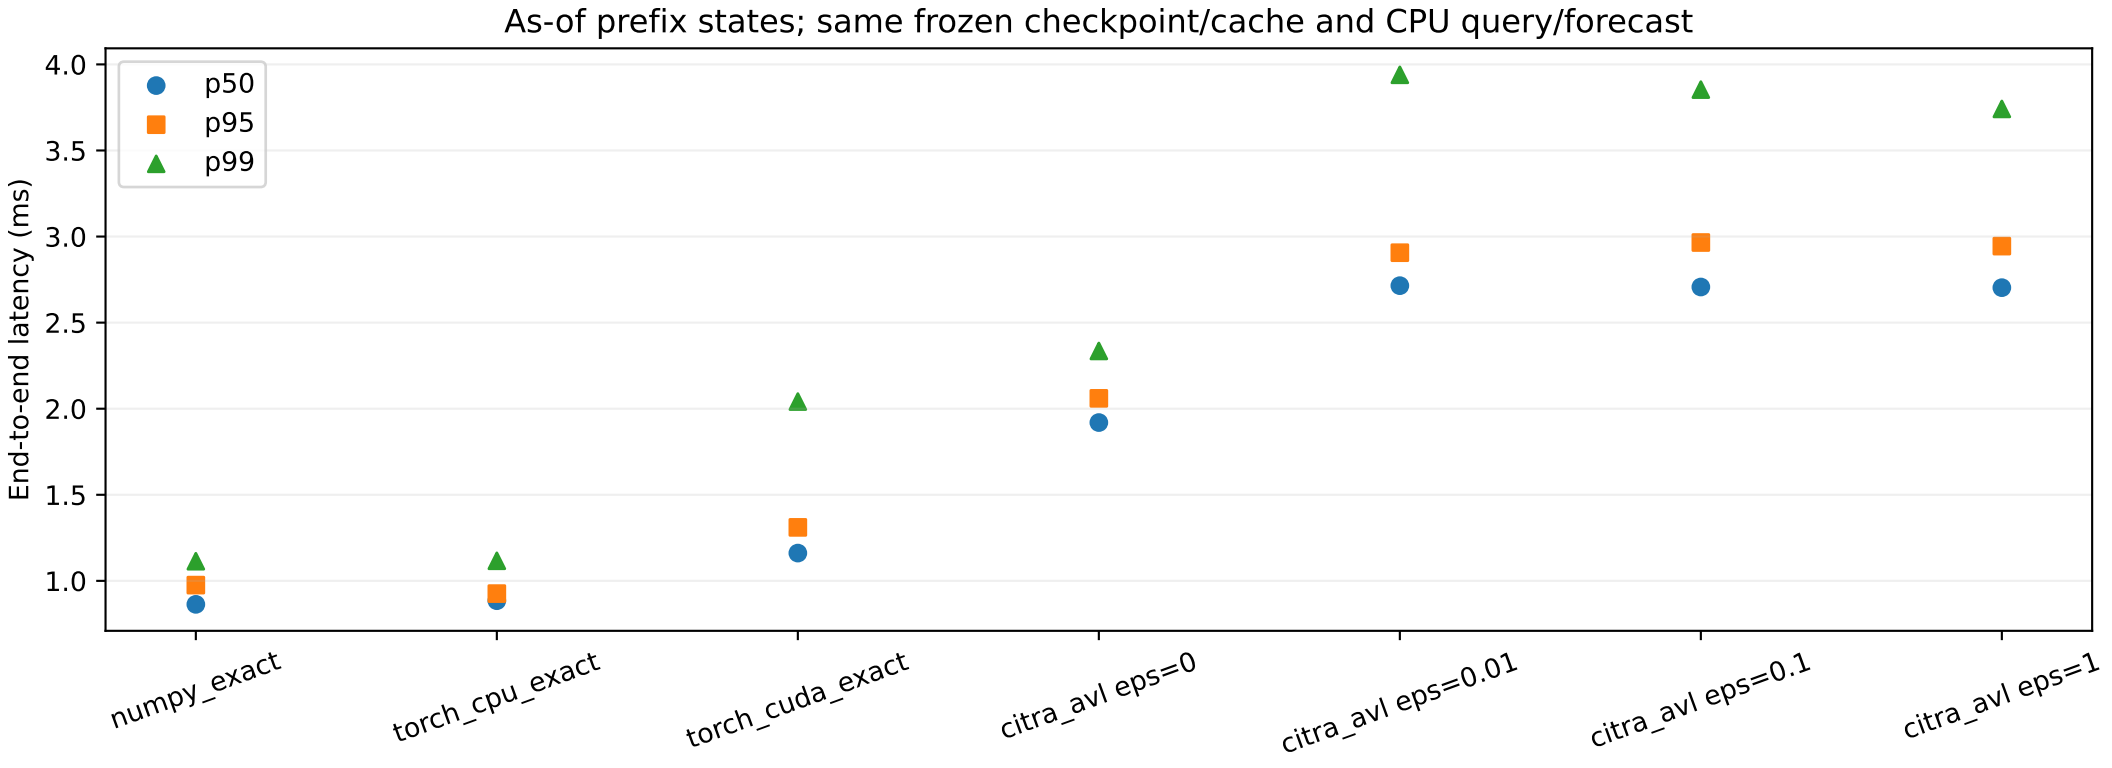
As-of prefix states; same frozen checkpoint/cache and CPU query/forecast
End-to-end latency (ms)
4.0
3.5
3.0
2.5
2.0
1.5
1.0
p50
p95
p99
numpy_exact
torch_cpu_exact
torch_cuda_exact
citra_avl eps=0
citra_avl eps=0.01
citra_avl eps=0.1
citra_avl eps=1